

Class 7: Machine Learning 1

Cynthia Perez (A16393492)


Today we will start our multi-part exploration of some key machine learning methods. We will begin with clustering- finding groupings in data, and then dimensionality reduction.

Clustering


Let's start with “k-means” clustering. The main function in base R for this is `kmeans()`

```
#Make up some data  
hist( rnorm(100000, mean=3) )
```

Histogram of `rnorm(1e+05, mean = 3)`



```
#add both rnorm values into one vector
tmp <- c(rnorm(30, -3), rnorm(30, +3))
# first 30 values start at -3 last 30 values are above 3
x <- cbind(x=tmp, y=rev(tmp))
plot(x)
```



Now that we have input data lets try `kmeans()`

```
km <- kmeans(x, centers=2)
km
```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:

	x	y
1	-2.936049	3.397500
2	3.397500	-2.936049

Clustering vector:

```
[1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2 2 2
[39] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```

```
[1] 50.95272 50.95272
      (between_SS / total_SS =  92.2 %)
```


```
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
[6] "betweenss"    "size"         "iter"         "ifault"
```

km\$size

km\$cluster

km\$centers


```
#color by cluster using `km$cluster` vector
plot(x, col=km$cluster)
#color the cluster centers
points(km$centers, col="green", pch=15, cex=3)
```



Q. Run `kmeans()` again and cluster in 4 groups and plot the results

```
# kmeans of vector x with 4 groups
km4 <- kmeans(x, centers = 4)

#plot km4
plot(x, col=km4$cluster)
```



Hierarchical Clustering

This form of clustering aims to reveal the structure in your data by progressively grouping points into smaller number of clusters.

The main function in base R for this is `hclust()`. This function does not take our input data directly but want a “distance matrix” that details how (dis)similar all our input points are to each other.

```
# `dist()` measures distance pairwise from point to point
hc <- hclust(dist(x))
hc
```

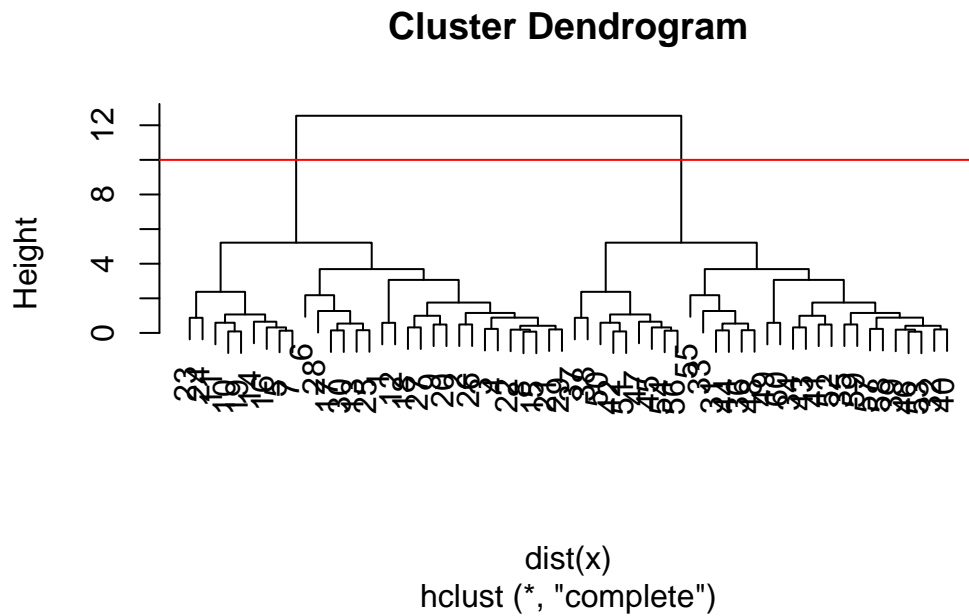
Call:

```
hclust(d = dist(x))
```

```
Cluster method   : complete
Distance         : euclidean
Number of objects: 60
```

The print out above is not very useful (unlike that from kmeans) but there is a useful `plot()` method.

```
plot(hc)
abline(h=10, col="red")
```




To get my main result (my cluster membership vector) I need to “cut” my tree using `cutree()`

```
#tree cut at height 10 creates 2 groups
grps <- cutree(hc, h=10)
grps
```

```
[1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2 2
[39] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```

Plot `hc` and color by groups (`grps`)

```
plot(x, col=grps)
```



##PCA of UK Food Data

```
#read the input file
url <- "https://tinyurl.com/UK-foods"
#use row.names to remove column name (x)
x <- read.csv(url, row.names=1)
x
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139
Fresh_potatoes	720	874	566	1033
Fresh_Veg	253	265	171	143
Other_Veg	488	570	418	355
Processed_potatoes	198	203	220	187
Processed_Veg	360	365	337	334
Fresh_fruit	1102	1137	957	674
Cereals	1472	1582	1462	1494

Beverages	57	73	53	47
Soft_drinks	1374	1256	1572	1506
Alcoholic_drinks	375	475	458	135
Confectionery	54	64	62	41


Q. How many rows and columns are in in data frame x?

```
dim(x)
```

```
[1] 17  4
```


Use barplot to spot trends

```
#change beside to TRUE to unstack the graph
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```



We can alternatively use the “pairs” plot for small datasets:

```
#color by number of food rows
#pch= plotting character to visualize the points better
pairs(x, col=rainbow(nrow(x)), pch=16)
```

The pairs plot is useful for small data sets but it can be too much work to interpret and becomes even harder to read with larger data sets.

Use PCA instead with the function `prcomp()`

```
# need to transpose x to perform PCA on the food and not the countries (switch the rows and columns)
pca <- prcomp(t(x))
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	3.176e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

Take a look at what is in `pca`

```
attributes(pca)
```

```
$names
[1] "sdev"      "rotation" "center"    "scale"     "x"
```

```
$class  
[1] "prcomp"
```

x is what the data looks like on the new axis


```
pca$x
```

	PC1	PC2	PC3	PC4
England	-144.99315	-2.532999	105.768945	-4.894696e-14
Wales	-240.52915	-224.646925	-56.475555	5.700024e-13
Scotland	-91.86934	286.081786	-44.415495	-7.460785e-13
N.Ireland	477.39164	-58.901862	-4.877895	2.321303e-13

```
#PC1 captures the most variance (makes it more important)
```

A major PCA result visualization is a “PCA plot” (aka a score plot, biplot, PC1 vs PC2 plot, ordination plot)

```
mycols <- c("orange", "red", "blue", "green")  
plot(pca$x[,1], pca$x[,2], col=mycols, pch=16,  
      xlab="PC1", ylab="PC2")  
abline(h=0, col="gray")  
abline(v=0, col="gray")
```



```
#Ireland sticks out as the green point
```

Another important output from PCA is called the “loadings” vector or the “rotation” component- this tells us how much the original variables (the food in this case) contribute to the new PCs.

```
pca$rotation
```

	PC1	PC2	PC3	PC4
Cheese	-0.056955380	0.016012850	0.02394295	-0.694538519
Carcass_meat	0.047927628	0.013915823	0.06367111	0.489884628
Other_meat	-0.258916658	-0.015331138	-0.55384854	0.279023718
Fish	-0.084414983	-0.050754947	0.03906481	-0.008483145
Fats_and_oils	-0.005193623	-0.095388656	-0.12522257	0.076097502
Sugars	-0.037620983	-0.043021699	-0.03605745	0.034101334
Fresh_potatoes	0.401402060	-0.715017078	-0.20668248	-0.090972715
Fresh_Veg	-0.151849942	-0.144900268	0.21382237	-0.039901917
Other_Veg	-0.243593729	-0.225450923	-0.05332841	0.016719075
Processed_potatoes	-0.026886233	0.042850761	-0.07364902	0.030125166
Processed_Veg	-0.036488269	-0.045451802	0.05289191	-0.013969507
Fresh_fruit	-0.632640898	-0.177740743	0.40012865	0.184072217
Cereals	-0.047702858	-0.212599678	-0.35884921	0.191926714

Beverages	-0.026187756	-0.030560542	-0.04135860	0.004831876
Soft_drinks	0.232244140	0.555124311	-0.16942648	0.103508492
Alcoholic_drinks	-0.463968168	0.113536523	-0.49858320	-0.316290619
Confectionery	-0.029650201	0.005949921	-0.05232164	0.001847469

PCA is a super useful method for gaining some insight into high dimensional data that is difficult to interpret in other ways.